

- Student name: Michael Holthouser
- Student pace: flex
- Scheduled project review date/time:
- Instructor name: Abhineet Kulkarni
- Blog post URL:

## Introduction

---

SyriaTel has tasked me to provide prediction analysis on whether their customers will churn soon. To churn in its broadest sense according to wikipedia is, "A measure of the number of individuals or items moving out of a collective group over a specific period."

## Stakeholder:

- SyriaTel Communications

## Data:

The title of this dataset is called "Churn in Telecom's dataset" from [kaggle.com](https://www.kaggle.com/datasets/becksddf/churn-in-telecoms-dataset) (<https://www.kaggle.com/datasets/becksddf/churn-in-telecoms-dataset>)

- Number of records: 3333
- Number of columns: 20
- Target variable: **churn**

## Models:

- Baseline model: Logistic regression
- Model 2: Decision tree
- Model 3: Random forest

## Evaluation Metric:

I have elected to use **Recall** as my evaluation metric for this particular project. The recall score is true positive divided by the true positive plus the false negative. It is the measure of actual observations which are predicted correctly. I chose this metric because we want to capture as many positives as possible, and is the best metric to use when we have imbalanced data.

# Import Libraries

Firstly, we must import the necessary library packages for this project.

```
In [1]: 1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 %matplotlib inline
5 import seaborn as sns
6 import warnings
7 warnings.filterwarnings('ignore')
8
9 from sklearn.pipeline import Pipeline
10 from sklearn.model_selection import train_test_split, GridSearchCV
11 from sklearn.preprocessing import OneHotEncoder
12 from sklearn.linear_model import LogisticRegression
13 from sklearn.metrics import confusion_matrix
14 from sklearn.metrics import plot_confusion_matrix, classification_report
15 from imblearn.over_sampling import RandomOverSampler
16
17
18 from sklearn.metrics import precision_score, recall_score, accuracy_score
19
20
21 from sklearn.tree import DecisionTreeClassifier
22 from sklearn import tree
23 from sklearn.ensemble import RandomForestClassifier, BaggingClassifier
```

## Column Descriptions

- **state** , string. 2-letter code of the US state of customer residence
- **account\_length** , numerical. Number of months the customer has been with the current telco provider
- **area\_code** , string="area\_code\_AAA" where AAA = 3 digit area code.
- **international\_plan** , (yes/no). The customer has international plan.
- **voice\_mail\_plan** , (yes/no). The customer has voice mail plan.
- **number\_vmail\_messages** , numerical. Number of voice-mail messages.
- **total\_day\_minutes** , numerical. Total minutes of day calls.
- **total\_day\_calls** , numerical. Total minutes of day calls.
- **total\_day\_charge** , numerical. Total charge of day calls.
- **total\_eve\_minutes** , numerical. Total minutes of evening calls.
- **total\_eve\_calls** , numerical. Total number of evening calls.
- **total\_eve\_charge** , numerical. Total charge of evening calls.
- **total\_night\_minutes** , numerical. Total minutes of night calls.
- **total\_night\_calls** , numerical. Total number of night calls.
- **total\_night\_charge** , numerical. Total charge of night calls.
- **total\_intl\_minutes** , numerical. Total minutes of international calls.
- **total\_intl\_calls** , numerical. Total number of international calls.

- **total\_intl\_charge** , numerical. Total charge of international calls
- **number\_customer\_service\_calls** , numerical. Number of calls to customer service
- **churn** , (yes/no). Customer churn - target variable.

## Loading The Data

The next step is to extract the data and put in a pandas dataframe, and to print the first 5 rows to see if the data was imported correctly.

```
In [2]: 1 data = pd.read_csv('churn_telecom.csv')
        2 data.head()
```

Out[2]:

	state	account length	area code	phone number	international plan	voice mail plan	number vmail messages	total day minutes	total day calls	total day charge	...	total eve calls
0	KS	128	415	382-4657	no	yes	25	265.1	110	45.07	...	99
1	OH	107	415	371-7191	no	yes	26	161.6	123	27.47	...	103
2	NJ	137	415	358-1921	no	no	0	243.4	114	41.38	...	110
3	OH	84	408	375-9999	yes	no	0	299.4	71	50.90	...	88
4	OK	75	415	330-6626	yes	no	0	166.7	113	28.34	...	122

5 rows × 21 columns

## Exploring and Cleaning the data

Below I am exploring the data, and checking what the data types of each columns with the .info() function, the descriptive statistics of the data with the .describe() function, and finally I will check for any missing data using the .isna() function.

```
In [3]: 1 # inspect how many records there are, and the data types for each column
        2 data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3333 entries, 0 to 3332
Data columns (total 21 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   state                                3333 non-null   object
 1   account length                       3333 non-null   int64
 2   area code                           3333 non-null   int64
 3   phone number                        3333 non-null   object
 4   international plan                  3333 non-null   object
 5   voice mail plan                     3333 non-null   object
 6   number vmail messages               3333 non-null   int64
 7   total day minutes                   3333 non-null   float64
 8   total day calls                     3333 non-null   int64
 9   total day charge                     3333 non-null   float64
10   total eve minutes                   3333 non-null   float64
11   total eve calls                     3333 non-null   int64
12   total eve charge                     3333 non-null   float64
13   total night minutes                 3333 non-null   float64
14   total night calls                   3333 non-null   int64
15   total night charge                   3333 non-null   float64
16   total intl minutes                  3333 non-null   float64
17   total intl calls                    3333 non-null   int64
18   total intl charge                   3333 non-null   float64
19   customer service calls              3333 non-null   int64
20   churn                              3333 non-null   bool
dtypes: bool(1), float64(8), int64(8), object(4)
memory usage: 524.2+ KB
```

At first glance, it appears for the most part the columns that are numerical are the correct data type. However, there are a couple columns that are the object data type that need to be changed to something numerical. The "international plan" and the "voice mail plan" have entries in the form of "yes" or "no" need to be changed to 1 for "yes" and 0 for "no". Likewise the "churn" column needs to be changed into a different data type and will have entries changed to numerical entries such as 1 for "True" and 0 for "False".

```
In [4]: 1 # Descriptive statistics on each column of the dataset
        2 data.describe()
```

Out[4]:

	account length	area code	number vmail messages	total day minutes	total day calls	total day charge	total eve minutes
<b>count</b>	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000
<b>mean</b>	101.064806	437.182418	8.099010	179.775098	100.435644	30.562307	200.980348
<b>std</b>	39.822106	42.371290	13.688365	54.467389	20.069084	9.259435	50.713844
<b>min</b>	1.000000	408.000000	0.000000	0.000000	0.000000	0.000000	0.000000
<b>25%</b>	74.000000	408.000000	0.000000	143.700000	87.000000	24.430000	166.600000
<b>50%</b>	101.000000	415.000000	0.000000	179.400000	101.000000	30.500000	201.400000
<b>75%</b>	127.000000	510.000000	20.000000	216.400000	114.000000	36.790000	235.300000
<b>max</b>	243.000000	510.000000	51.000000	350.800000	165.000000	59.640000	363.700000

The data looks normal without and large outliers glooming.

```
In [5]: 1 # Inspect the dataset to see if there is any missing data
        2 data.isna().sum()
```

```
Out[5]: state                0
account length              0
area code                  0
phone number               0
international plan         0
voice mail plan            0
number vmail messages     0
total day minutes          0
total day calls            0
total day charge           0
total eve minutes          0
total eve calls            0
total eve charge           0
total night minutes        0
total night calls          0
total night charge         0
total intl minutes         0
total intl calls           0
total intl charge          0
customer service calls     0
churn                      0
dtype: int64
```

Great! there is no missing values to take care of. However, I will next add an underscore "\_" as that is conventional in the python language.

```
In [6]: 1 data.columns = data.columns.map(lambda col: col.replace(' ', '_'))
        2 print(data.columns)

Index(['state', 'account_length', 'area_code', 'phone_number',
       'international_plan', 'voice_mail_plan', 'number_vmail_messages',
       'total_day_minutes', 'total_day_calls', 'total_day_charge',
       'total_eve_minutes', 'total_eve_calls', 'total_eve_charge',
       'total_night_minutes', 'total_night_calls', 'total_night_charge',
       'total_intl_minutes', 'total_intl_calls', 'total_intl_charge',
       'customer_service_calls', 'churn'],
      dtype='object')
```

Next I shall check the unique values of the dataset to see if there are any place holders for missing values.

```

In [7]: 1 # inspect unique values of columns to identify potention errors or null
        2 for col in data.columns:
        3     print(f"{col} vals: {data[col].unique()} \n")

state vals: ['KS' 'OH' 'NJ' 'OK' 'AL' 'MA' 'MO' 'LA' 'WV' 'IN' 'RI' 'IA'
'MT' 'NY'
'ID' 'VT' 'VA' 'TX' 'FL' 'CO' 'AZ' 'SC' 'NE' 'WY' 'HI' 'IL' 'NH' 'GA'
'AK' 'MD' 'AR' 'WI' 'OR' 'MI' 'DE' 'UT' 'CA' 'MN' 'SD' 'NC' 'WA' 'NM'
'NV' 'DC' 'KY' 'ME' 'MS' 'TN' 'PA' 'CT' 'ND']

account_length vals: [128 107 137 84 75 118 121 147 117 141 65 74 16
8 95 62 161 85 93
76 73 77 130 111 132 174 57 54 20 49 142 172 12 72 36 78 136
149 98 135 34 160 64 59 119 97 52 60 10 96 87 81 68 125 116
38 40 43 113 126 150 138 162 90 50 82 144 46 70 55 106 94 155
80 104 99 120 108 122 157 103 63 112 41 193 61 92 131 163 91 127
110 140 83 145 56 151 139 6 115 146 185 148 32 25 179 67 19 170
164 51 208 53 105 66 86 35 88 123 45 100 215 22 33 114 24 101
143 48 71 167 89 199 166 158 196 209 16 39 173 129 44 79 31 124
37 159 194 154 21 133 224 58 11 109 102 165 18 30 176 47 190 152
26 69 186 171 28 153 169 13 27 3 42 189 156 134 243 23 1 205
200 5 9 178 181 182 217 177 210 29 180 2 17 7 212 232 192 195
197 225 184 191 201 15 183 202 8 175 4 188 204 221]

area_code vals: [415 408 510]

phone_number vals: ['382-4657' '371-7191' '358-1921' ... '328-8230' '364
-6381' '400-4344']

international_plan vals: ['no' 'yes']

voice_mail_plan vals: ['yes' 'no']

number_vmail_messages vals: [25 26 0 24 37 27 33 39 30 41 28 34 46 29 3
5 21 32 42 36 22 23 43 31 38
40 48 18 17 45 16 20 14 19 51 15 11 12 47 8 44 49 4 10 13 50 9]

total_day_minutes vals: [265.1 161.6 243.4 ... 321.1 231.1 180.8]

total_day_calls vals: [110 123 114 71 113 98 88 79 97 84 137 127
96 70 67 139 66 90
117 89 112 103 86 76 115 73 109 95 105 121 118 94 80 128 64 106
102 85 82 77 120 133 135 108 57 83 129 91 92 74 93 101 146 72
99 104 125 61 100 87 131 65 124 119 52 68 107 47 116 151 126 122
111 145 78 136 140 148 81 55 69 158 134 130 63 53 75 141 163 59
132 138 54 58 62 144 143 147 36 40 150 56 51 165 30 48 60 42
0 45 160 149 152 142 156 35 49 157 44]

total_day_charge vals: [45.07 27.47 41.38 ... 54.59 39.29 30.74]

total_eve_minutes vals: [197.4 195.5 121.2 ... 153.4 288.8 265.9]

total_eve_calls vals: [99 103 110 88 122 101 108 94 80 111 83 148
71 75 76 97 90 65
93 121 102 72 112 100 84 109 63 107 115 119 116 92 85 98 118 74
117 58 96 66 67 62 77 164 126 142 64 104 79 95 86 105 81 113
106 59 48 82 87 123 114 140 128 60 78 125 91 46 138 129 89 133

```

```

136 57 135 139 51 70 151 137 134 73 152 168 68 120 69 127 132 143
61 124 42 54 131 52 149 56 37 130 49 146 147 55 12 50 157 155
45 144 36 156 53 141 44 153 154 150 43 0 145 159 170]

```

```
total_eve_charge vals: [16.78 16.62 10.3 ... 13.04 24.55 22.6 ]
```

```
total_night_minutes vals: [244.7 254.4 162.6 ... 280.9 120.1 279.1]
```

```

total_night_calls vals: [ 91 103 104 89 121 118 96 90 97 111 94 128
115 99 75 108 74 133
64 78 105 68 102 148 98 116 71 109 107 135 92 86 127 79 87 129
57 77 95 54 106 53 67 139 60 100 61 73 113 76 119 88 84 62
137 72 142 114 126 122 81 123 117 82 80 120 130 134 59 112 132 110
101 150 69 131 83 93 124 136 125 66 143 58 55 85 56 70 46 42
152 44 145 50 153 49 175 63 138 154 140 141 146 65 51 151 158 155
157 147 144 149 166 52 33 156 38 36 48 164]

```

```

total_night_charge vals: [11.01 11.45 7.32 8.86 8.41 9.18 9.57 9.5
3 9.71 14.69 9.4 8.82
6.35 8.65 9.14 7.23 4.02 5.83 7.46 8.68 9.43 8.18 8.53 10.67
11.28 8.22 4.59 8.17 8.04 11.27 11.08 13.2 12.61 9.61 6.88 5.82
10.25 4.58 8.47 8.45 5.5 14.02 8.03 11.94 7.34 6.06 10.9 6.44
3.18 10.66 11.21 12.73 10.28 12.16 6.34 8.15 5.84 8.52 7.5 7.48
6.21 11.95 7.15 9.63 7.1 6.91 6.69 13.29 11.46 7.76 6.86 8.16
12.15 7.79 7.99 10.29 10.08 12.53 7.91 10.02 8.61 14.54 8.21 9.09
4.93 11.39 11.88 5.75 7.83 8.59 7.52 12.38 7.21 5.81 8.1 11.04
11.19 8.55 8.42 9.76 9.87 10.86 5.36 10.03 11.15 9.51 6.22 2.59
7.65 6.45 9. 6.4 9.94 5.08 10.23 11.36 6.97 10.16 7.88 11.91
6.61 11.55 11.76 9.27 9.29 11.12 10.69 8.8 11.85 7.14 8.71 11.42
4.94 9.02 11.22 4.97 9.15 5.45 7.27 12.91 7.75 13.46 6.32 12.13
11.97 6.93 11.66 7.42 6.19 11.41 10.33 10.65 11.92 4.77 4.38 7.41
12.1 7.69 8.78 9.36 9.05 12.7 6.16 6.05 10.85 8.93 3.48 10.4
5.05 10.71 9.37 6.75 8.12 11.77 11.49 11.06 11.25 11.03 10.82 8.91
8.57 8.09 10.05 11.7 10.17 8.74 5.51 11.11 3.29 10.13 6.8 8.49
9.55 11.02 9.91 7.84 10.62 9.97 3.44 7.35 9.79 8.89 8.14 6.94
10.49 10.57 10.2 6.29 8.79 10.04 12.41 15.97 9.1 11.78 12.75 11.07
12.56 8.63 8.02 10.42 8.7 9.98 7.62 8.33 6.59 13.12 10.46 6.63
8.32 9.04 9.28 10.76 9.64 11.44 6.48 10.81 12.66 11.34 8.75 13.05
11.48 14.04 13.47 5.63 6.6 9.72 11.68 6.41 9.32 12.95 13.37 9.62
6.03 8.25 8.26 11.96 9.9 9.23 5.58 7.22 6.64 12.29 12.93 11.32
6.85 8.88 7.03 8.48 3.59 5.86 6.23 7.61 7.66 13.63 7.9 11.82
7.47 6.08 8.4 5.74 10.94 10.35 10.68 4.34 8.73 5.14 8.24 9.99
13.93 8.64 11.43 5.79 9.2 10.14 12.11 7.53 12.46 8.46 8.95 9.84
10.8 11.23 10.15 9.21 14.46 6.67 12.83 9.66 9.59 10.48 8.36 4.84
10.54 8.39 7.43 9.06 8.94 11.13 8.87 8.5 7.6 10.73 9.56 10.77
7.73 3.47 11.86 8.11 9.78 9.42 9.65 7. 7.39 9.88 6.56 5.92
6.95 15.71 8.06 4.86 7.8 8.58 10.06 5.21 6.92 6.15 13.49 9.38
12.62 12.26 8.19 11.65 11.62 10.83 7.92 7.33 13.01 13.26 12.22 11.58
5.97 10.99 8.38 9.17 8.08 5.71 3.41 12.63 11.79 12.96 7.64 6.58
10.84 10.22 6.52 5.55 7.63 5.11 5.89 10.78 3.05 11.89 8.97 10.44
10.5 9.35 5.66 11.09 9.83 5.44 10.11 6.39 11.93 8.62 12.06 6.02
8.85 5.25 8.66 6.73 10.21 11.59 13.87 7.77 10.39 5.54 6.62 13.33
6.24 12.59 6.3 6.79 8.28 9.03 8.07 5.52 12.14 10.59 7.54 7.67
5.47 8.81 8.51 13.45 8.77 6.43 12.01 12.08 7.07 6.51 6.84 9.48
13.78 11.54 11.67 8.13 10.79 7.13 4.72 4.64 8.96 13.03 6.07 3.51
6.83 6.12 9.31 9.58 4.68 5.32 9.26 11.52 9.11 10.55 11.47 9.3
13.82 8.44 5.77 10.96 11.74 8.9 10.47 7.85 10.92 4.74 9.74 10.43

```



```

9.96 10.18 9.54 7.89 12.36 8.54 10.07 9.46 7.3 11.16 9.16 10.19
5.99 10.88 5.8 7.19 4.55 8.31 8.01 14.43 8.3 14.3 6.53 8.2
11.31 13. 6.42 4.24 7.44 7.51 13.1 9.49 6.14 8.76 6.65 10.56
6.72 8.29 12.09 5.39 2.96 7.59 7.24 4.28 9.7 8.83 13.3 11.37
9.33 5.01 3.26 11.71 8.43 9.68 15.56 9.8 3.61 6.96 11.61 12.81
10.87 13.84 5.03 5.17 2.03 10.34 9.34 7.95 10.09 9.95 7.11 9.22
6.13 11.05 9.89 9.39 14.06 10.26 13.31 15.43 16.39 6.27 10.64 11.5
12.48 8.27 13.53 10.36 12.24 8.69 10.52 9.07 11.51 9.25 8.72 6.78
8.6 11.84 5.78 5.85 12.3 5.76 12.07 9.6 8.84 12.39 10.1 9.73
2.85 6.66 2.45 5.28 11.73 10.75 7.74 6.76 6. 7.58 13.69 7.93
7.68 9.75 4.96 5.49 11.83 7.18 9.19 7.7 7.25 10.74 4.27 13.8
9.12 4.75 7.78 11.63 7.55 2.25 9.45 9.86 7.71 4.95 7.4 11.17
11.33 6.82 13.7 1.97 10.89 12.77 10.31 5.23 5.27 9.41 6.09 10.61
7.29 4.23 7.57 3.67 12.69 14.5 5.95 7.87 5.96 5.94 12.23 4.9
12.33 6.89 9.67 12.68 12.87 3.7 6.04 13.13 15.74 11.87 4.7 4.67
7.05 5.42 4.09 5.73 9.47 8.05 6.87 3.71 15.86 7.49 11.69 6.46
10.45 12.9 5.41 11.26 1.04 6.49 6.37 12.21 6.77 12.65 7.86 9.44
4.3 7.38 5.02 10.63 2.86 17.19 8.67 8.37 6.9 10.93 10.38 7.36
10.27 10.95 6.11 4.45 11.9 15.01 12.84 7.45 6.98 11.72 7.56 11.38
10. 4.42 9.81 5.56 6.01 10.12 12.4 16.99 5.68 11.64 3.78 7.82
9.85 13.74 12.71 10.98 10.01 9.52 7.31 8.35 11.35 9.5 14.03 3.2
7.72 13.22 10.7 8.99 10.6 13.02 9.77 12.58 12.35 12.2 11.4 13.91
3.57 14.65 12.28 5.13 10.72 12.86 14. 7.12 12.17 4.71 6.28 8.
7.01 5.91 5.2 12. 12.02 12.88 7.28 5.4 12.04 5.24 10.3 10.41
13.41 12.72 9.08 7.08 13.5 5.35 12.45 5.3 10.32 5.15 12.67 5.22
5.57 3.94 4.41 13.27 10.24 4.25 12.89 5.72 12.5 11.29 3.25 11.53
9.82 7.26 4.1 10.37 4.98 6.74 12.52 14.56 8.34 3.82 3.86 13.97
11.57 6.5 13.58 14.32 13.75 11.14 14.18 9.13 4.46 4.83 9.69 14.13
7.16 7.98 13.66 14.78 11.2 9.93 11. 5.29 9.92 4.29 11.1 10.51
12.49 4.04 12.94 7.09 6.71 7.94 5.31 5.98 7.2 14.82 13.21 12.32
10.58 4.92 6.2 4.47 11.98 6.18 7.81 4.54 5.37 7.17 5.33 14.1
5.7 12.18 8.98 5.1 14.67 13.95 16.55 11.18 4.44 4.73 2.55 6.31
2.43 9.24 7.37 13.42 12.42 11.8 14.45 2.89 13.23 12.6 13.18 12.19
14.81 6.55 11.3 12.27 13.98 8.23 15.49 6.47 13.48 13.59 13.25 17.77
13.9 3.97 11.56 14.08 13.6 6.26 4.61 12.76 15.76 6.38 3.6 12.8
5.9 7.97 5. 10.97 5.88 12.34 12.03 14.97 15.06 12.85 6.54 11.24
12.64 7.06 5.38 13.14 3.99 3.32 4.51 4.12 3.93 2.4 11.75 4.03
15.85 6.81 14.25 14.09 16.42 6.7 12.74 2.76 12.12 6.99 6.68 11.81
7.96 5.06 13.16 2.13 13.17 5.12 5.65 12.37 10.53]

```

```

total_intl_minutes vals: [10. 13.7 12.2 6.6 10.1 6.3 7.5 7.1 8.7 1
1.2 12.7 9.1 12.3 13.1
5.4 13.8 8.1 13. 10.6 5.7 9.5 7.7 10.3 15.5 14.7 11.1 14.2 12.6
11.8 8.3 14.5 10.5 9.4 14.6 9.2 3.5 8.5 13.2 7.4 8.8 11. 7.8
6.8 11.4 9.3 9.7 10.2 8. 5.8 12.1 12. 11.6 8.2 6.2 7.3 6.1
11.7 15. 9.8 12.4 8.6 10.9 13.9 8.9 7.9 5.3 4.4 12.5 11.3 9.
9.6 13.3 20. 7.2 6.4 14.1 14.3 6.9 11.5 15.8 12.8 16.2 0. 11.9
9.9 8.4 10.8 13.4 10.7 17.6 4.7 2.7 13.5 12.9 14.4 10.4 6.7 15.4
4.5 6.5 15.6 5.9 18.9 7.6 5. 7. 14. 18. 16. 14.8 3.7 2.
4.8 15.3 6. 13.6 17.2 17.5 5.6 18.2 3.6 16.5 4.6 5.1 4.1 16.3
14.9 16.4 16.7 1.3 15.2 15.1 15.9 5.5 16.1 4. 16.9 5.2 4.2 15.7
17. 3.9 3.8 2.2 17.1 4.9 17.9 17.3 18.4 17.8 4.3 2.9 3.1 3.3
2.6 3.4 1.1 18.3 16.6 2.1 2.4 2.5]

```

```

total_intl_calls vals: [ 3 5 7 6 4 2 9 19 1 10 15 8 11 0 12 13
18 14 16 20 17]

```

```
total_intl_charge vals: [2.7  3.7  3.29 1.78 2.73 1.7  2.03 1.92 2.35 3.
02 3.43 2.46 3.32 3.54
 1.46 3.73 2.19 3.51 2.86 1.54 2.57 2.08 2.78 4.19 3.97 3.   3.83 3.4
 3.19 2.24 3.92 2.84 2.54 3.94 2.48 0.95 2.3  3.56 2.   2.38 2.97 2.11
 1.84 3.08 2.51 2.62 2.75 2.16 1.57 3.27 3.24 3.13 2.21 1.67 1.97 1.65
 3.16 4.05 2.65 3.35 2.32 2.94 3.75 2.4  2.13 1.43 1.19 3.38 3.05 2.43
 2.59 3.59 5.4  1.94 1.73 3.81 3.86 1.86 3.11 4.27 3.46 4.37 0.   3.21
 2.67 2.27 2.92 3.62 2.89 4.75 1.27 0.73 3.65 3.48 3.89 2.81 1.81 4.16
 1.22 1.76 4.21 1.59 5.1  2.05 1.35 1.89 3.78 4.86 4.32 4.   1.   0.54
 1.3  4.13 1.62 3.67 4.64 4.73 1.51 4.91 0.97 4.46 1.24 1.38 1.11 4.4
 4.02 4.43 4.51 0.35 4.1  4.08 4.29 1.49 4.35 1.08 4.56 1.4  1.13 4.24
 4.59 1.05 1.03 0.59 4.62 1.32 4.83 4.67 4.97 4.81 1.16 0.78 0.84 0.89
 0.7  0.92 0.3  4.94 4.48 0.57 0.65 0.68]

customer_service_calls vals: [1 0 2 3 4 5 7 9 6 8]

churn vals: [False  True]
```

Some things I noticed:

- **States** look good.
- There are only 3 area codes
- "**international\_plan**", "**voice\_mail\_plan**" have yes/no values and will need to be changed to a 1 and 0.
- Phone number can probably be dropped from the dataset because, it shouldn't be a reason why customer is choosing to churn.
- The target variable "**churn**" has a boolean value, and needs to be changed to a 1 and 0.

## Map columns: international plan, voice mail plan, and churn

```
In [8]: 1 data['international_plan'] = data['international_plan'].map({'no': 0, '
2 data['voice_mail_plan'] = data['voice_mail_plan'].map({'no': 0, 'yes':
3 data['churn'] = data['churn'].map({False: 0, True: 1})
4 data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3333 entries, 0 to 3332
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   state                                3333 non-null   object
1   account_length                      3333 non-null   int64
2   area_code                          3333 non-null   int64
3   phone_number                       3333 non-null   object
4   international_plan                 3333 non-null   int64
5   voice_mail_plan                   3333 non-null   int64
6   number_vmail_messages             3333 non-null   int64
7   total_day_minutes                 3333 non-null   float64
8   total_day_calls                   3333 non-null   int64
9   total_day_charge                  3333 non-null   float64
10  total_eve_minutes                 3333 non-null   float64
11  total_eve_calls                   3333 non-null   int64
12  total_eve_charge                  3333 non-null   float64
13  total_night_minutes              3333 non-null   float64
14  total_night_calls                3333 non-null   int64
15  total_night_charge               3333 non-null   float64
16  total_intl_minutes               3333 non-null   float64
17  total_intl_calls                 3333 non-null   int64
18  total_intl_charge                3333 non-null   float64
19  customer_service_calls           3333 non-null   int64
20  churn                           3333 non-null   int64
dtypes: float64(8), int64(11), object(2)
memory usage: 546.9+ KB
```

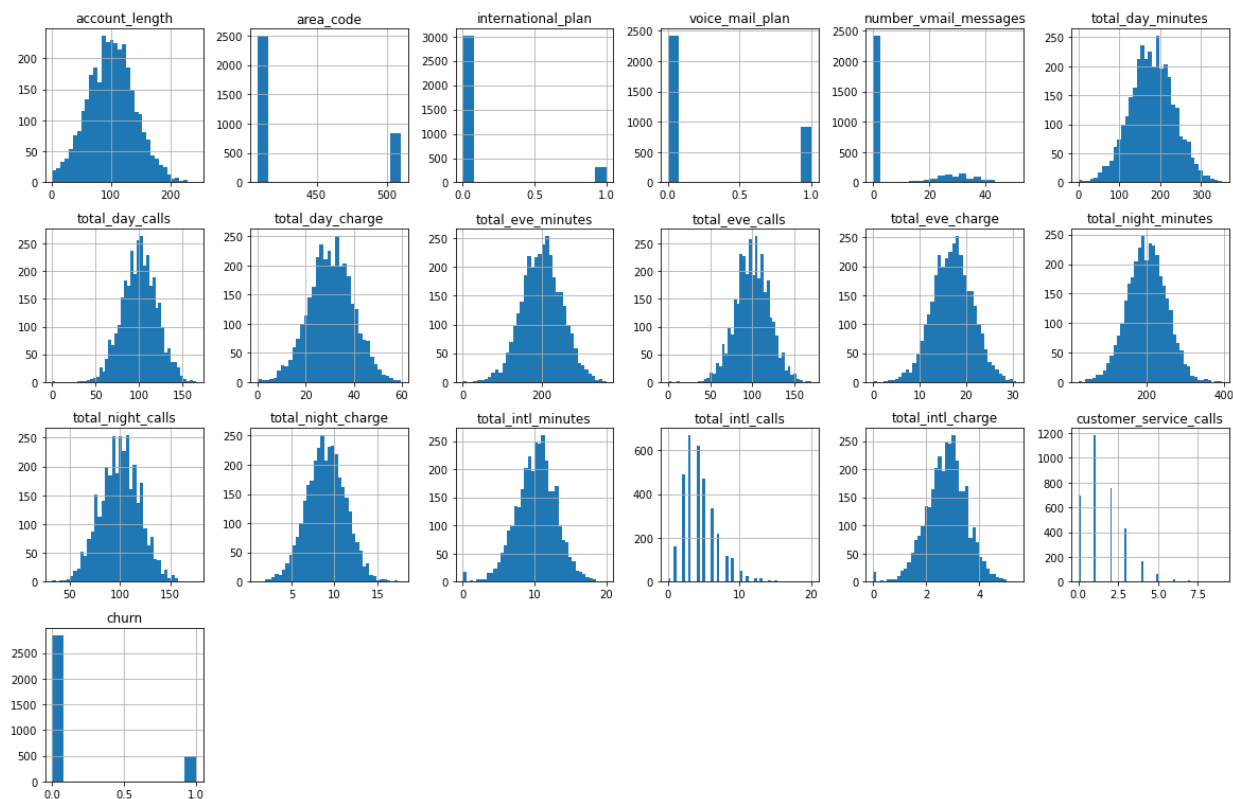
Now that we have dealt with the object data types, we can delete the phone number column, since it will not effect our models.

## Drop phone number column

```
In [9]: 1 # Drop the phone_number column from the dataset
2 data.drop("phone_number", axis=1, inplace=True)
```

## Check the distribution of the data

```
In [10]: 1 data.hist(bins = 'auto', layout = (6,6), figsize = (20,20))
        2 plt.show()
```



It looks like our data is normally distributed.

## Exploration and Visualization of Churn Data

Next I am curious and want visualize what the number of customers have churned and their percentages.

```
In [11]: 1 print("Churn Counts")
          2 print(data["churn"].value_counts())
          3 print()
          4 print("Percentages")
          5 print(data["churn"].value_counts(normalize=True))
```

Churn Counts

0 2850

1 483

Name: churn, dtype: int64

Percentages

0 0.855086

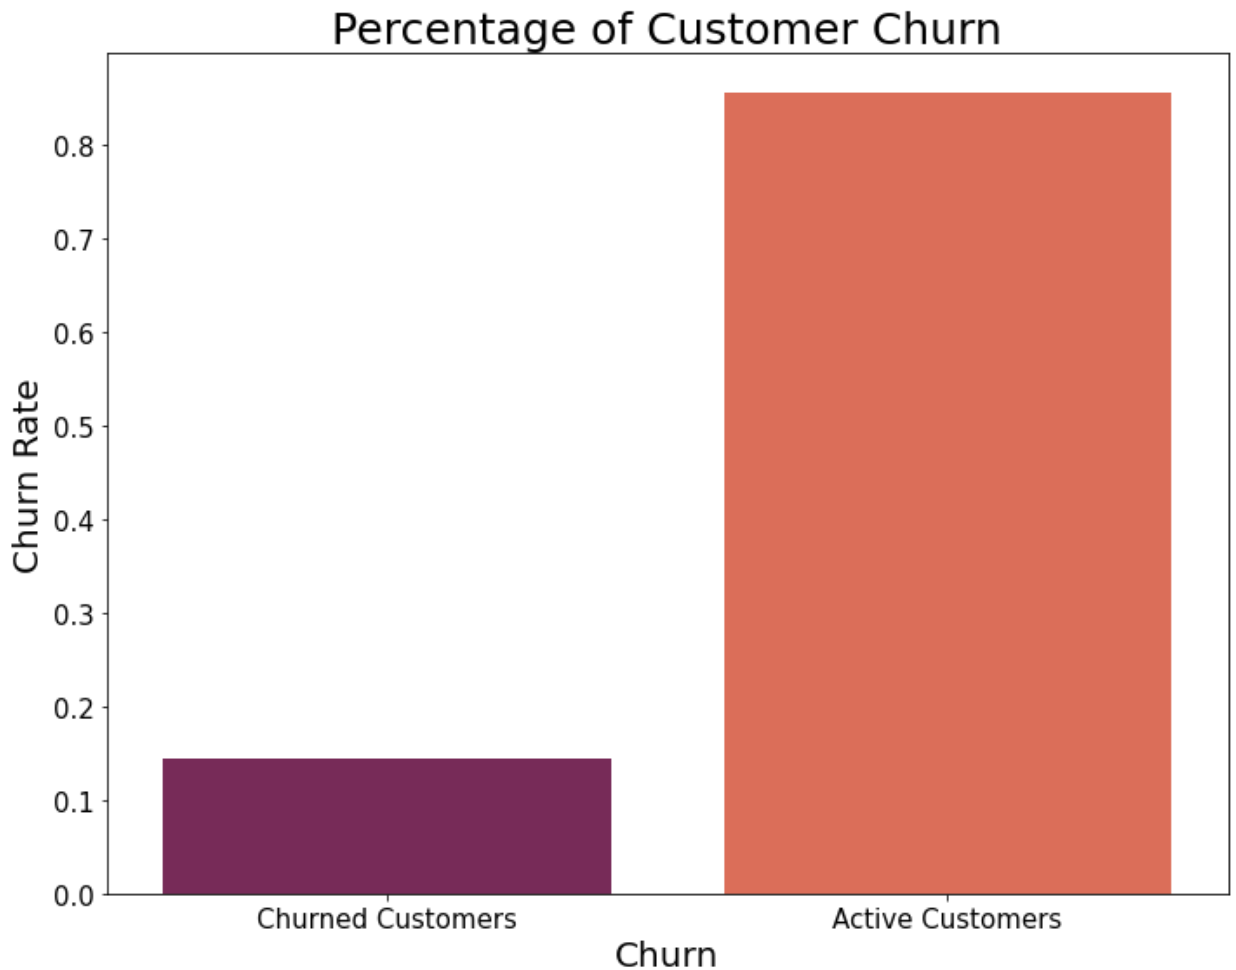
1 0.144914

Name: churn, dtype: float64

Of the 3333 customers from the data set **14.5%** have terminated their service with SyriaTel. I am curious if a certain area code has a greater number of churns over the other.

## Visualization of percentage of customers that have churned

```
In [12]: 1 # Percentages of current customers vs customer churn
2 churn_per = data["churn"].value_counts(normalize=True)
3
4 # Plot of percentages
5 fig, ax = plt.subplots(figsize = (10, 8))
6 sns.barplot(x = [0, 1], y = data["churn"], palette="rocket", data = churn_per)
7 plt.title('Percentage of Customer Churn', fontsize = 25)
8 ax.tick_params(axis='both', labelsize=15)
9 plt.xlabel('Churn', fontsize=20)
10 plt.ylabel('Churn Rate', fontsize=20)
11 ax.set_xticklabels(['Churned Customers', 'Active Customers'])
12 plt.tight_layout()
```



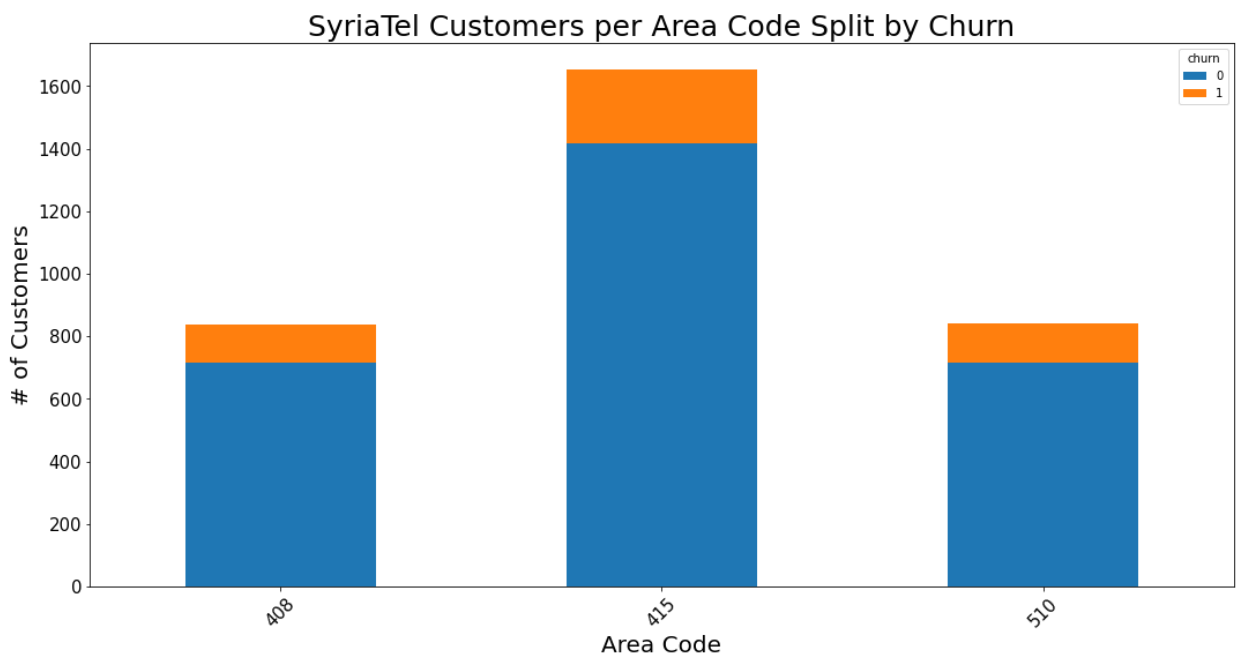
This is a visualization to show the churn rate of the SyriaTel customers. As shown above, the churn rate for this particular dataset was **14.5%**

## Visualization of number of customers churned per area code

Below I wanted to see if there was any pattern with the churn rate versus the area codes provided in the dataset.

```
In [13]: 1 # percentage of churn by area code
2 print(data.groupby(["area_code"])[ 'churn' ].mean())
3
4 fig, ax = plt.subplots(figsize = (15, 8))
5 data.groupby(['area_code', 'churn']).size().unstack().plot(kind='bar',
6 plt.title('SyriaTel Customers per Area Code Split by Churn', fontsize =
7 ax.tick_params(axis = 'both', labelsize = 15)
8 plt.xlabel('Area Code', fontsize = 20)
9 plt.ylabel('# of Customers', fontsize = 20)
10 plt.tight_layout()
11
12 ##rotate x-axis to a 45 degree angle
13 for label in ax.xaxis.get_ticklabels():
14     label.set_rotation(45)
```

```
area_code
408    0.145585
415    0.142598
510    0.148810
Name: churn, dtype: float64
```



After further investigation, it is clear that the 415 area code has more customers than the 408 or 510 area codes. However, all three area codes have around the same **churn rate**. Since there is no clear pattern, I believe it is safe to delete the area code column from the dataset as well.

- **churn rate** – is the rate at which customers or clients are leaving a company within a specific period of time.

## Drop area code column

Since area codes had no influence on the churn rate, I think it is safe to remove the column from our dataset.

```
In [14]: 1 # Drop the area_code column from the dataset
          2 data.drop('area_code', axis=1, inplace=True)
```

```
In [15]: 1 # Check to see if phone_number and area_code columns have been removed
          2 data.columns
```

```
Out[15]: Index(['state', 'account_length', 'international_plan', 'voice_mail_plan',
               'number_vmail_messages', 'total_day_minutes', 'total_day_calls',
               'total_day_charge', 'total_eve_minutes', 'total_eve_calls',
               'total_eve_charge', 'total_night_minutes', 'total_night_calls',
               'total_night_charge', 'total_intl_minutes', 'total_intl_calls',
               'total_intl_charge', 'customer_service_calls', 'churn'],
              dtype='object')
```

## EDA continued: Correlation between churn and other features

Below I used the `.corr()` function to see what features had the highest correlation with the target variable, churn.

```
In [16]: 1 # Correlation with Churn
          2 data.corr().churn.sort_values(ascending=False)
```

```
Out[16]: churn                1.000000
international_plan          0.259852
customer_service_calls      0.208750
total_day_minutes           0.205151
total_day_charge             0.205151
total_eve_minutes           0.092796
total_eve_charge             0.092786
total_intl_charge            0.068259
total_intl_minutes           0.068239
total_night_charge           0.035496
total_night_minutes          0.035493
total_day_calls              0.018459
account_length               0.016541
total_eve_calls              0.009233
total_night_calls            0.006141
total_intl_calls             -0.052844
number_vmail_messages        -0.089728
voice_mail_plan              -0.102148
Name: churn, dtype: float64
```

international\_plan, customer\_service\_calls, and total\_day\_charge/total\_day\_minutes have the highest correlation with churn. This will be later confirmed when I do the feature importance investigation.



## Average churn rate for international plan holders


```
In [17]: 1 inter_plan_churn = pd.DataFrame(data.groupby(['international_plan'])['churn']
          2 inter_plan_churn
```

Out[17]:

	churn
international_plan	
0	0.114950
1	0.424149

It appears 42% customers with an international plan with SyriaTel, end up churning. On a business stand point, this may be a worthwhile topic to further investigate.

```
In [18]: 1 # International plan bar plot
          2 fig, ax = plt.subplots(figsize=(13,7))
          3 sns.barplot(data=data, x=data['international_plan'], y='churn', palette
          4 plt.title('Percentage Churn of International Plan Holders', fontsize=25)
          5 plt.xlabel('International Plan Yes/No', fontsize=20)
          6 plt.ylabel('Churn Rate', fontsize=20)
          7 ax.set_xticklabels(['No International Plan', 'Has International Plan'],
          8 plt.tight_layout;
```



For customer service calls, you would imagine the more calls a customer must make to customer service, the likely they are to be unhappy with their phone service. But how many calls on average does it take to increase the likely hood for a customer to churn. Let's take a look.


## Average churn rate for number of customer service calls

```
In [19]: 1 cust_serv_calls = pd.DataFrame(data.groupby(['customer_service_calls']))
          2 cust_serv_calls
```

Out[19]:

churn	
customer_service_calls	
0	0.131994
1	0.103302
2	0.114625
3	0.102564
4	0.457831
5	0.606061
6	0.636364
7	0.555556
8	0.500000
9	1.000000

```
In [20]: 1 # Bar plot for customer service calls
          2 fig, ax = plt.subplots(figsize=(13,7))
          3 sns.barplot(data=data, x=data['customer_service_calls'], y='churn',pale
          4 plt.title('Churn Rate in Relation to Customer Service Calls', fontsize=
          5 plt.ylabel('Churn Rate', fontsize=20)
          6 plt.xlabel('Number of Calls', fontsize=20)
          7 plt.tight_layout();
```



From the graph above, it is evident that when a customer has to call customer service four times, the likely hood of a customer to churn significantly increases. When a customer needs to call a maximum 9 times, the churn rate reaches 100%. Looking at this in a business perspective, new

strategies must be discussed to handle unhappy customers when they are calling customer service by the fourth time.

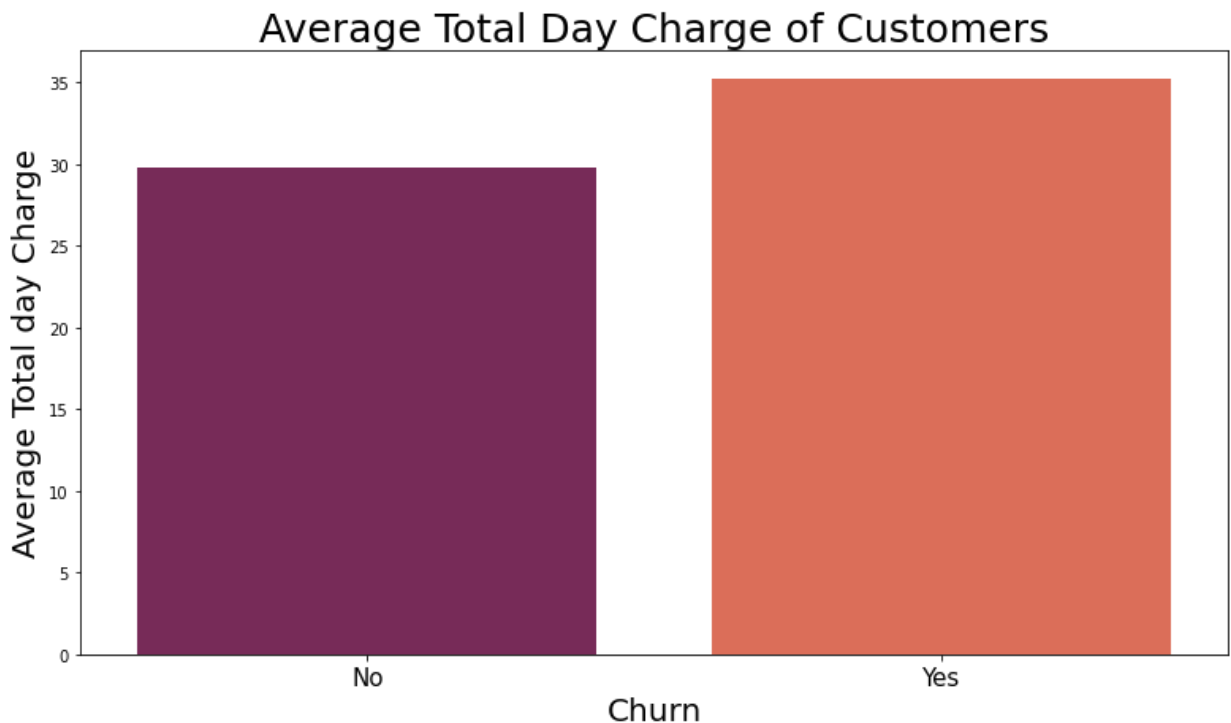
## Average total day charge for churned and current customers

```
In [21]: 1 tdc = data.groupby(data['churn'])\
2         ['total_day_charge'].mean().sort_values(ascending=False).reset_index()
3 tdc.head()
```

Out[21]:

	churn	total_day_charge
0	1	35.175921
1	0	29.780421

```
In [22]: 1 # total_day_charge plan bar plot
2 fig, ax = plt.subplots(figsize=(13,7))
3 sns.barplot(data=tdc, x='churn', y='total_day_charge', palette="rocket")
4 plt.title('Average Total Day Charge of Customers', fontsize=25)
5 plt.xlabel('Churn', fontsize=20)
6 plt.ylabel('Average Total day Charge', fontsize=20)
7 ax.set_xticklabels(['No', 'Yes'], fontsize=15)
8 plt.tight_layout;
```



Graph above shows the average day charge a SyriaTel customer has, and when they are on average churning. Per the data, when a customer is around the 35.18 mark for their average day charge, they are more likely to churn. 29.78 and lower is the price the company should strive to be around in order to keep their customers from churning.

## Create dummy variables for state column

Before modeling, there is one column that must be dealt with. The state column is categorical, and needs to be converted to dummy variables and added to the dataframe.

```
In [23]: 1 state_dum = pd.get_dummies(data['state'], drop_first=True)
```

```
In [24]: 1 data_final = data.drop('state', axis=1)
```

```
In [25]: 1 data_final = pd.concat([data_final, state_dum], axis=1)
2 data_final.head()
```

Out[25]:

	account_length	international_plan	voice_mail_plan	number_vmail_messages	total_day_minutes	t
0	128	0	1	25	265.1	
1	107	0	1	26	161.6	
2	137	0	0	0	243.4	
3	84	1	0	0	299.4	
4	75	1	0	0	166.7	

5 rows × 68 columns

## Prepare Data for Modeling

### Create X, y variables

```
1 Below I am creating the target variable y, and the independant
   variables X.
```

```
In [26]: 1 # Create X
2 X = data_final.drop('churn', axis=1)
3
4 # Create y
5 y = data_final['churn']
```

### Train, Test, Split

below I am splitting the data into a training set and test set

```
In [27]: 1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

## Baseline Model

## Logistic Regression

```
In [28]: 1 # Initialize logistic regression
2 logreg = LogisticRegression(solver='liblinear', random_state=42)
3
4 # fit the model
5 logreg.fit(X_train, y_train)
```

Out[28]: LogisticRegression(random\_state=42, solver='liblinear')

### Get Predictions

```
In [29]: 1 y_hat_train = logreg.predict(X_train)
2 y_hat_test = logreg.predict(X_test)
```

### Classification report of the training data

```
In [30]: 1 display(confusion_matrix(y_train, y_hat_train))
2 print(classification_report(y_train, y_hat_train))
```

```
array([[2099,  42],
       [ 278,  80]])
```

	precision	recall	f1-score	support
0	0.88	0.98	0.93	2141
1	0.66	0.22	0.33	358
accuracy			0.87	2499
macro avg	0.77	0.60	0.63	2499
weighted avg	0.85	0.87	0.84	2499

```
In [31]: 1 print("Training Accuracy for Logistic Regression: {:.4}%".format(accura
```

Training Accuracy for Logistic Regression: 87.19%

After running the first classification report on the training data, our baseline model had a recall score of 22%. This is not very good and leaves a lot of room for improvement.

### Check for imbalance

```
In [32]: 1 # check for imbalance of the training data
2 print(y_train.value_counts())
3 print('\n')
4 print(y_test.value_counts())
```

```
0    2141
1     358
Name: churn, dtype: int64
```

```
0     709
1     125
Name: churn, dtype: int64
```

It looks like we are dealing with some imbalance here. We could use SMOTE to deal with imbalance for future work.

### Classification report of the testing data

```
In [33]: 1 display(confusion_matrix(y_test, y_hat_test))
2 print(classification_report(y_test, y_hat_test))
```

```
array([[688, 21],
       [106, 19]])
```

	precision	recall	f1-score	support
0	0.87	0.97	0.92	709
1	0.47	0.15	0.23	125
accuracy			0.85	834
macro avg	0.67	0.56	0.57	834
weighted avg	0.81	0.85	0.81	834

```
In [34]: 1 y_test.value_counts()
```

```
Out[34]: 0     709
1     125
Name: churn, dtype: int64
```

```
In [35]: 1 print("Test Accuracy for Logistic Regression: {:.4}%".format(accuracy_s
Test Accuracy for Logistic Regression: 84.77%
```

```
1 Our test set had an even lower score than our training set with a
  recall score of only 15%. On a positive note, we did manage to lower
  our false negative score. Let's try a different model.
```

### Results:

- Training data recall score: 22%
- Test data recall score: 15%

## Model 2: Decision Tree

### Train the Decision Tree

```
In [36]: 1 clf = DecisionTreeClassifier(max_depth = 5, min_samples_leaf = 2, min_s
          2
          3 clf.fit(X_train, y_train)
```

```
Out[36]: DecisionTreeClassifier(max_depth=5, min_samples_leaf=2, min_samples_split
        =5,
        random_state=42)
```

Decision trees require some pruning to become more accurate. For this model I used min\_samples\_split, and min\_samples\_leaf.

### Evaluate the Predictive Performance

```
In [37]: 1 # Make prediction for test data
          2 y_pred = clf.predict(X_test)
          3 y_pred_train = clf.predict(X_train)
```

### Display classification report for training data

```
In [38]: 1 #print confusion matrix and classification report
          2 display(confusion_matrix(y_train, y_pred_train))
          3 print(classification_report(y_train, y_pred_train))
```

```
array([[2129,   12],
       [ 101,  257]])
```

	precision	recall	f1-score	support
0	0.95	0.99	0.97	2141
1	0.96	0.72	0.82	358
accuracy			0.95	2499
macro avg	0.96	0.86	0.90	2499
weighted avg	0.95	0.95	0.95	2499

```
In [39]: 1 print("Training Accuracy for Decision Tree: {:.4}%".format(accuracy_sco
Training Accuracy for Decision Tree: 95.48%
```

We have a much better recall score on our training data using the decision tree model. We improved from our baseline model from 22% previously, to 72% with the decision tree model.

**Display classification report for test data**

```
In [40]: 1 #print confusion matrix and classification report
2 display(confusion_matrix(y_test, y_pred))
3 print(classification_report(y_test, y_pred))
```

```
array([[699, 10],
       [ 43, 82]])
```

	precision	recall	f1-score	support
0	0.94	0.99	0.96	709
1	0.89	0.66	0.76	125
accuracy			0.94	834
macro avg	0.92	0.82	0.86	834
weighted avg	0.93	0.94	0.93	834

```
In [41]: 1 print("Test Accuracy for Decision Tree: {:.4}%".format(accuracy_score(y
```

```
Test Accuracy for Decision Tree: 93.65%
```

We have gotten a slightly lower recall score on our test data, but still scored much better than our baseline model. Also, our false negative number was reduced even further to 43.

**Results:**

- Training data recall score: 72%
- Test data recall score: 66%

## Model 3: Random Forest Classifier

**Create X and y variables**

```
In [42]: 1 # assign X and y variables
2 X = data_final.drop('churn', axis =1)
3 y = data_final['churn']
```

**Split the data into training and testing sets**

```
In [43]: 1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2
```

**Fit a random forest model**



```
In [44]: 1 #Instantiate and fit a RandomForestClassifier
2 rf = RandomForestClassifier(n_estimators=10, max_depth=5, random_state=
3 rf.fit(X_train, y_train)
```

```
Out[44]: RandomForestClassifier(max_depth=5, n_estimators=10, random_state=42)
```

## Get prediction and build classification report

### Get predictions for training and test data

```
In [45]: 1 # get predictions
2 y_pred = rf.predict(X_test)
3 y_pred_train = rf.predict(X_train)
```

### Print confusion matrix and classification report for training data

```
In [46]: 1 #build a confusion matrix and classification report
2 display(confusion_matrix(y_train, y_pred_train))
3 print(classification_report(y_train, y_pred_train))
```

```
array([[2138,    3],
       [ 251,  107]])
```

	precision	recall	f1-score	support
0	0.89	1.00	0.94	2141
1	0.97	0.30	0.46	358
accuracy			0.90	2499
macro avg	0.93	0.65	0.70	2499
weighted avg	0.91	0.90	0.87	2499

```
In [47]: 1 print("Training Accuracy for Random Forest: {:.4}%".format(accuracy_sco
```

```
Training Accuracy for Random Forest: 89.84%
```

On our training data, the random forest model did not perform as well as the decision tree producing a recall score of 30%.

### Print confusion matrix and classification report for test data

```
In [48]: 1 #build a confusion matrix and classification report
2 display(confusion_matrix(y_test, y_pred))
3 print(classification_report(y_test, y_pred))
```

```
array([[705,  4],
       [ 96, 29]])
```

	precision	recall	f1-score	support
0	0.88	0.99	0.93	709
1	0.88	0.23	0.37	125
accuracy			0.88	834
macro avg	0.88	0.61	0.65	834
weighted avg	0.88	0.88	0.85	834

```
In [49]: 1 print("Test Accuracy for Random Forest: {:.4}%".format(accuracy_score(y
```

```
Test Accuracy for Random Forest: 88.01%
```

On our test data, random forest forest again produced another low recall score of 23%. The false negative number was not improved from the decision tree model.

### Results:

- Training data recall score: 30%
- Test data recall score: 23%

## Tuning the model with GridSearchCV

I am now performing GridSearchCV on the random forest model, to see what hypertuning should be taken place in order to get the best performing random forest model.

```
In [50]: 1 param_grid = {'criterion':['gini','entropy'],
2                     'max_depth':[3,6,9,12],
3                     'min_samples_leaf': [5,10,12],
4                     'n_estimators': [10, 12, 15, 20]
5                     }
```

### Fit the gridsearch

```
In [51]: 1 rfc = RandomForestClassifier(n_estimators=10, max_depth=5, random_state=42)
2 grid_rfc = GridSearchCV(rfc, param_grid, cv=5, scoring='accuracy')
3 grid_rfc.fit(X_train, y_train)
```

```
Out[51]: GridSearchCV(cv=5,
                    estimator=RandomForestClassifier(max_depth=5, n_estimators=10,
                    random_state=42),
                    param_grid={'criterion': ['gini', 'entropy'],
                    'max_depth': [3, 6, 9, 12],
                    'min_samples_leaf': [5, 10, 12],
                    'n_estimators': [10, 12, 15, 20]},
                    scoring='accuracy')
```

### Display the gridsearch results

```
In [52]: 1 # Print the best parameters
2 grid_rfc.best_params_
```

```
Out[52]: {'criterion': 'gini',
          'max_depth': 12,
          'min_samples_leaf': 5,
          'n_estimators': 12}
```

With these recommendations, I will apply this hypertunings to the model to see if we get better results.

### Get predictions from gridsearch

```
In [53]: 1 # get predictions from the gridsearch
2 y_pred_grid = grid_rfc.predict(X_test)
3 y_pred_train_grid = grid_rfc.predict(X_train)
```

### Print confusion matrix and classification report for training data

```
In [54]: 1 #print confusion matrix and classification report
2 display(confusion_matrix(y_train, y_pred_train_grid))
3 print(classification_report(y_train, y_pred_train_grid))
```

```
array([[2140,    1],
       [ 105,   253]])
```

	precision	recall	f1-score	support
0	0.95	1.00	0.98	2141
1	1.00	0.71	0.83	358
accuracy			0.96	2499
macro avg	0.97	0.85	0.90	2499
weighted avg	0.96	0.96	0.95	2499

```
In [55]: 1 print("Training Accuracy for Random Forest Classifier: {:.4}%".format(a
```

```
Training Accuracy for Random Forest Classifier: 95.76%
```

Great! We saw an improvement in our recall score from 30% to 71% on our training data after using gridSearch

### Print confusion matrix and classification report for test data

```
In [56]: 1 #print confusion matrix and classification report
2 display(confusion_matrix(y_test, y_pred_grid))
3 print(classification_report(y_test, y_pred_grid))
```

```
array([[705,    4],
       [ 57,   68]])
```

	precision	recall	f1-score	support
0	0.93	0.99	0.96	709
1	0.94	0.54	0.69	125
accuracy			0.93	834
macro avg	0.93	0.77	0.82	834
weighted avg	0.93	0.93	0.92	834

```
In [57]: 1 print("Test Accuracy for Random Forest Classifier: {:.4}%".format(accur
```

```
Test Accuracy for Random Forest Classifier: 92.69%
```

The recall score took a little dip when running it on our test data, resulting a recall score of 54%.  
The false negative number did decrease to 57.

### Results:

- Training data recall score: 71%

- Test data recall score: 54%

## Feature Importance


Next we will examine how important each feature ended up being in our final model. In machine learning, feature selection is an important step. More features equals more complex models that take longer to train and are harder to interpret.

```
In [58]: 1 # Feature importance
          2 rf.feature_importances_
```

```
Out[58]: array([1.22494433e-02, 1.62707771e-01, 4.76229153e-02, 2.38905778e-02,
                1.46961828e-01, 1.09306397e-02, 1.61786301e-01, 4.02865581e-02,
                2.94028749e-02, 3.18338451e-02, 1.00894253e-02, 1.22712100e-02,
                2.53034318e-02, 1.49793444e-02, 2.70418095e-02, 3.13637003e-02,
                1.59027372e-01, 1.81158198e-03, 1.61274014e-03, 0.00000000e+00,
                2.79599790e-03, 1.89470501e-03, 0.00000000e+00, 0.00000000e+00,
                5.84599959e-04, 1.72440166e-03, 0.00000000e+00, 0.00000000e+00,
                2.72819229e-04, 9.15745437e-04, 0.00000000e+00, 4.87626773e-04,
                5.26425267e-03, 3.84937889e-03, 0.00000000e+00, 0.00000000e+00,
                2.25138758e-03, 5.96763790e-03, 0.00000000e+00, 7.69026920e-04,
                0.00000000e+00, 0.00000000e+00, 4.58178301e-03, 0.00000000e+00,
                0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 3.99977346e-03,
                7.21375938e-04, 1.16223313e-03, 1.96181241e-03, 2.05621859e-04,
                3.03082977e-03, 0.00000000e+00, 1.97705363e-03, 1.49307190e-03,
                0.00000000e+00, 0.00000000e+00, 1.36076802e-04, 0.00000000e+00,
                0.00000000e+00, 0.00000000e+00, 1.66817014e-03, 0.00000000e+00,
                0.00000000e+00, 1.11004588e-03, 1.20263870e-06])
```

This array full of numbers isn't very helpful. Let's plot the data to see if the important features become more clear.

```
In [59]: 1 def plot_features_importances(model):
2         n_features = X_test.shape[1]
3         plt.figure(figsize=(15,20))
4         plt.barh(range(n_features), model.feature_importances_, align='center')
5         plt.yticks(np.arange(n_features), X_test.columns.values, fontsize = 25)
6         plt.xlabel('Feature Importance', fontsize = 25)
7         plt.ylabel('Feature', fontsize = 25)
8         plt.title('Final Model Feature Importance', fontsize = 25)
9         plt.tight_layout()
10
11 plot_features_importances(rf)
```



we can see from this feature importance graph that there are three features that the model is weighing more heavily, with little to no weight given to the states.

- **total\_day\_charge**
- **customer\_service\_calls**
- **international\_plan**

## Conclusion

### Logistic Regression:

- Recall Score (Training): 22%
- Recall Score (Test): 15%

### Decision Tree:

- Recall Score (Training): 72%
- Recall Score (Test): 66%

### Random Forest:

- Recall Score (Training): 30%
- Recall Score (Test): 23%

### Random Forest with GridSearchCV:

- Recall Score (Training): 71%
- Recall Score (Test): 54%

From our findings, I can conclude that the decision tree model was the best testing model having the highest recall score on it's training data as well as it's test data.

## Recommendations:

- 42% percent of the customers that have churned had international plans. Further discussion and investigations should be taken place to formulate a plan to retain these customers.
- Customers that have called customer service at least 4 times have a significantly increased chance of churning. Managers must come up with new training techniques to help customer service representatives assist these disgruntled customers.
- Investigate ways to retain customers that have an average total day charge of 35 dollars. Possibly creating more incentives and added perks to their phone plans could sway these customers from terminating their contracts.